



ASHTAVADHANI (अष्टवधानी)

Multi-Agent Collaborative Video Synthesis Engine

Author: **Sanyam Asopa** (sanyam4747@gmail.com) | GitHub: github.com/sanyam47/ashtavadhani

1. Executive Summary

Ashtavadhani is an edge-first, multi-agent collaborative video synthesis and zero-shot viral edit cloning engine. Inspired by the classical Indian intellectual tradition of *Ashtavadhana*—performing eight complex cognitive tasks concurrently without context loss—Ashtavadhani orchestrates a council of specialized AI agents and neural vision models to automate high-end short-form video editing directly on modern AI PCs.

By harnessing **Qualcomm Snapdragon X Elite NPU hardware acceleration** via the Qualcomm AI Hub (qai-hub), Ashtavadhani delivers sub-millisecond neural execution, 100% NPU compute offload, and complete on-device data sovereignty—eliminating recurring cloud AI subscription fees and ensuring creator media never leaks to remote servers.

2. Problem Statement

- **The Creative Production Bottleneck:** Creating viral 30-second reels featuring kinetic typography, beat synchronization, and mask cutouts takes 4–8 hours of manual labor in professional NLE software (Premiere Pro, After Effects).
- **Cloud AI Latency & Skyrocketing Costs:** Existing generative video platforms rely on centralized cloud GPUs, imposing long rendering queues, high bandwidth consumption, and recurring \$50–\$100/month subscriptions.
- **Severe Creator Privacy Risks:** Raw personal footage, brand campaigns, and unreleased family media are exposed to third-party cloud servers.
- **Underutilized NPU Silicon:** Modern Copilot+ AI PCs feature 45 TOPS NPUs (like the Qualcomm Snapdragon X Elite), yet 99% of creative video editing tools remain stuck on legacy CPU/GPU software rendering.

3. The Solution: Dual-Mode Architecture

Ashtavadhani features a unified dual-mode architecture tailored for both narrative storytelling and precision viral replication:

Mode 1: The 8-Agent Council

1. **Manager Agent:** Scope & timeline choreography
2. **Vision Agent:** Semantic clip inspection & quality scoring
3. **Ref Analyzer:** Viral rhythm & pacing pattern extraction
4. **Music Agent:** FFT audio wave profiling & beat drop sync
5. **SFX Agent:** Contextual sound hits, whooshes & impacts
6. **Transitions Agent:** Dynamic directional whip-pans & zooms
7. **Motion Graphics:** Vibe grading (Phonk, Teal/Orange LUTs)
8. **Caption Typographer:** Word-level kinetic subtitles

Mode 2: 9-Stage Neural Pipeline

1. **Snapdragon NPU:** Hardware acceleration verification
2. **Blueprint Analyzer:** Keyframe rhythm & typography timing
3. **BiRefNet Matting:** Sub-pixel portrait subject cutout
4. **ISNet Masking:** Temporal subject tracking across frames
5. **Model Dispatcher:** Visual entropy background routing
6. **LaMa Inpainting:** Fourier Convolution plate synthesis
7. **Anatomical Aligner:** Eye-locked landmark alignment
8. **3D Depth Composer:** Screen-light typography depth
9. **Audio Sync & Mux:** Lossless beat-locked AAC/MP4 delivery

4. Qualcomm Snapdragon Hardware Acceleration & Benchmarks

Ashtavadhani is deeply integrated with the **Qualcomm AI Hub SDK (qai-hub 0.55.0)**. Models were compiled using `--target_runtime precompiled_qnn_onnx` and validated on physical Snapdragon X Elite hardware.

100% NPU COMPUTE OFFLOAD Zero CPU/GPU inference bottleneck	0.52 ms INFERENCE LATENCY Benchmarked on Snapdragon X Elite CRD	36.2 MB PEAK MEMORY FOOTPRINT Ultra-low footprint leaves RAM free
---	--	--

- **Resilient 3-Tier Fallback Chain:** Snapdragon NPU (Primary) → CUDA GPU (Secondary) → Multithreaded CPU (Universal Fallback). Runs universally while delivering exponential speedups on Snapdragon PCs.
- **All-Day Battery Efficiency:** Hexagon NPU offloading extends laptop battery life up to 3–4x compared to power-hungry discrete GPU renderers.

5. Key Algorithmic Breakthroughs

- **Visual Entropy Model Dispatcher:** Dynamically analyzes Margin Hue Std-Dev and Laplacian Edge Density to route between Fast Plate Averaging, Fourier LaMa Inpainting, and ProPainter.
- **Smart Anatomical Landmark Eye-Lock:** Eliminates uncanny distortions by anchoring user subjects via reference eye-lines and shoulder-width proportions.
- **Screen-Light 3D Typography Compositor:** Blends kinetic typography using Screen Light transfer modes behind or in front of creators with zero hardcoded fonts.
- **Dynamic Reference Memory Cache:** Serializes keyframe segmentations, cutting re-render cycles from minutes to seconds.

6. Market Potential & Creator Value

- **Global Reach:** Democratizes viral video editing for 200M+ digital creators, indie filmmakers, and marketing agencies.
- **Economic Advantage:** 10x faster turnaround; \$0/month in recurring server API fees; 100% sovereign media privacy on local hardware.